



The Geography of Open Source Software: Evidence from GitHub

Johannes Wachs^{a,b,*}, Mariusz Nitecki^a, William Schueller^{b,c}, Axel Polleres^{a,b}

^a Vienna University of Economics and Business

^b Complexity Science Hub Vienna

^c Medical University of Vienna

ARTICLE INFO

Keywords:

Geography
Open source software
GitHub
Innovation
cluster policy

ABSTRACT

Open Source Software (OSS) plays an important role in the digital economy. Yet although software production is amenable to remote collaboration and its outputs are digital, software development seems to cluster geographically in places like Silicon Valley, London, or Berlin. And while OSS activity creates positive externalities which accrue locally through knowledge spillovers and information effects, up-to-date data on the geographic distribution of open source developers is limited. This presents a significant blindspot for policy-makers, who often promote OSS at the national level as a cost-saving tool for public sector institutions. We address this gap by geolocating more than half a million active contributors to GitHub in early 2021 at various spatial scales. Compared to results from 2010, we find a significant increase in the share of developers based in Asia, Latin America and Eastern Europe, suggesting a more even spread of OSS developers globally. Within countries, however, we find significant concentration in regions, exceeding the concentration of high-tech employment. Social and economic development indicators predict at most half of regional variation in OSS activity in the EU, suggesting that clusters have idiosyncratic roots. We argue for localized policies to support networks of OSS developers in cities and regions.

Introduction

The importance of software, both as a ubiquitous complement to other activities in the modern economy and as a key sector in its own right, is widely acknowledged (Andreessen (2011); Nagle (2019b)). Software is also an especially global industry, in part because its end products are easily shared and distributed through the Web. Yet the notion that the software industry might transcend geographic constraints is inconsistent with anecdotal observations that its most coveted jobs and leading firms tend to cluster in particular places like Silicon Valley, London, or Berlin (Grier (2015)). These places have an out-sized influence on the art and practice of software engineering (Takhetev (2012)), which is increasingly shaping our economy and society (Wagner et al. (2021)). They also benefit from significant positive local externalities (Nagle et al. (2020)). Despite these observations, researchers and policymakers have limited information about the extent of spatial clustering of software development and the socioeconomic conditions in these hotspots.

We address this gap by studying the geography of open source software (OSS) developers. Within the vast ecosystem of software, OSS plays

a distinguished role (Eghbal (2020)), and is sometimes described as the infrastructure of the digital society. One recent estimate is that OSS contributes about 60 to 95 billion Euros annually to Eurozone GDP, and that a 10% increase in OSS activity would lead to roughly 600 additional ICT startups in the EU (Blind et al. (2021)). Mirroring trends in closed-source software, OSS contributions are thought to be intensely concentrated in space, despite the low cost of distributing software, the development of technologies for remote collaboration, and its inherently open nature. Research from 2010 estimated that over 7% of global OSS activity at the time took place in or around San Francisco (Takhetev and Hiltz (2010)). As OSS activity is known to impact local firm productivity (Nagle (2018, 2019b)) and rates of technology entrepreneurship (Wright et al. (2020)), its geographic clustering likely effects economic growth and inequality, hence should be of interest to policy-makers (Nagle (2019a)). Policy interventions to date have sought to foster OSS for example through procurement regulations or trade and industrial policy (Blind et al. (2021)). These policies are usually motivated by the potential cost-savings of OSS adoption by public sector institutions, instead of the positive externalities that OSS production creates for local economies.

* Corresponding author.

E-mail address: johannes.wachs@wu.ac.at (J. Wachs).

<https://doi.org/10.1016/j.techfore.2022.121478>

Received 7 July 2021; Received in revised form 16 December 2021; Accepted 6 January 2022

Available online 15 January 2022

0040-1625/© 2022 The Author(s). Published by Elsevier Inc. This is an open access article under the CC BY license (<http://creativecommons.org/licenses/by/4.0/>).

Geographic disparities may also influence who participates in OSS, and understanding them better can inform us about root causes of diversity issues in software (Albusays et al. (2021); May et al. (2019); Prana et al. (2021)) and its role in innovation (Dahlander et al. (2021)). Though OSS is thought to be a highly decentralized activity in the global economy, previous work has shown how core places in the software industry have incredible influence over the way software is written around the world (Takhteyev (2012)), for example by setting the norms and practices of software development used around the world. And though some breakthrough OSS innovations come from places such as Japan (Ruby), Finland (Linux), and South Africa (Ubuntu), they are often refined and popularized through central places like Chicago (where Ruby on Rails was created), the Pacific Northwest (where the creator of Linux now lives and works), and London (where Ubuntu is maintained) (Takhteyev and Hiltz (2010)).

So while the geographic distribution of OSS contributors likely has important localized effects on our society and economy, we lack an up-to-date geographic mapping of OSS activity. Previous work was either carried out over ten years ago (Gonzalez-Barahona et al. (2008); Takhteyev and Hiltz (2010)), focuses on data at the country-level (Nagle (2019a); Wright et al. (2020)), and/or on a subset of projects (Prana et al. (2021)). We suggest that there is need for a fresh look at the geographic distribution of OSS developers, including regional data. This gap is becoming more relevant as researchers take a greater interest in the influence of algorithms and software on society (Wagner et al. (2021)). If algorithms do shape the fabric of our society, the hypothesis that software is created by a distinct group of people living in a few particular places merits testing. At the same time, software engineering researchers have gathered massive datasets on OS contributions (Fry et al. (2020); Gousios and Spinellis (2012); Pietri et al. (2019)), presenting opportunities to study this question in greater depth.

In this work we implement a pipeline to geolocate highly active OSS developers on GitHub, cleaning and sharing the resulting data at both national and subnational levels. In contrast to studies from over ten years ago, we find a slightly more even global distribution of OSS activity, with marked growth among Asian, Latin American, and Eastern European countries. Within countries, however, we find that OSS remains highly concentrated in particular regions. OSS activity is even more spatially concentrated than the university educated population and workers in high-tech sectors. To better understand the structural conditions of geographic OSS hubs, we use a regression framework to relate OSS activity with social, political, and economic features of countries. No single feature can predict most of the variance observed between countries. A spatial econometric study of the EU regions replicates this finding at a more granular scale. This suggests that there are many necessary but not sufficient conditions for OSS clusters to develop in a place, with important implications for policymakers seeking to encourage OSS activity.

The rest of the paper is organized as follows. We first discuss related works on the importance of OSS in the digital economy both in general and locally. We then introduce our data collection process, describe how we geolocate developers, and share access to data on counts of active developers in various geographic units. We then analyze the distribution of developers at both the country and regional levels. We conclude with a discussion of limitations, potential uses for this dataset, and ideas for future work in this area.

Literature Review

In this section we review the literature on why industries cluster geographically in general, and reflect on ways in which the software industry and OSS in particular are unique in this context. We then survey emerging evidence on the local impact of OSS developers. Finally, we review the motivations for and results of policy interventions involving OSS. Research gaps between these concepts motivate our subsequent mapping of OSS activity and our regression analysis of correlated socio-

economic factors.

Geographic Clusters and the case of Open Source Software

It has long been observed that highly specialized activity at the frontier of knowledge clusters geographically (Krugman (1991)). Skilled people living and working near to one another interact and come up with new ideas, attracting more people with their own ideas (Watney (2020)). Network, information and agglomeration effects give these places comparative advantages that make up for pathologies of boom towns such as high rents. Many examples have been studied in the literature, for example bio-tech clusters in the US and China (Su and Hung (2009)), car manufacturing in Detroit (Klepper (2010)), and the shoe industry in southern Italy (Boschma and Ter Wal (2007)). The latter study in particular highlights that co-location alone is insufficient to create the virtuous effects associated with successful clusters that are thought to accelerate innovation, entrepreneurship, and productivity growth. In particular it is the social networks and cognitive proximities of co-located firms and people of a cluster that make it sink or swim (Frenken et al. (2015); Juhász (2021); Juhász and Lengyel (2018)). Silicon Valley itself bounced back from significant setbacks in the 1980s due to its regional networks of people and institutions (Saxenian (1990)). Despite the highly specific and seemingly irreplicable nature of individual clusters, they appear in a large variety of sectors.

The software industry is no exception. Indeed, even though software is easily shared and distributed, software creation has long been intensely geographically clustered (Bettencourt et al. (2002)). A 2015 study claimed that 40% of all jobs in the US software industry were clustered in merely nine cities (Grier (2015)) and Silicon Valley is a household name. As in other sectors, geographic proximity to customers, end-users, and other software developers improves developer productivity and software quality via spillover effects (Weterings and Boschma (2006)). Knowledge transfers between developers in the software industry often occur via informal networks, which complement more formal structures of collaboration and interaction (Trippi et al. (2009)).

Within the world of software nowadays, OSS activity is also thought to be highly localized, even if long distance collaborations are frequent and frictionless thanks to the sophisticated collaboration tools OSS developers use. A study of activity in the 2000s estimated that roughly one in ten contributions to OSS libraries on GitHub globally originated in the San Francisco Bay Area (Takhteyev and Hiltz (2010)). Several studies of OSS (Fackler and Laurensyeva (2020); Lima et al. (2014); Takhteyev and Hiltz (2010)) have shown that the likelihood of collaboration decreases exponentially with distance. Moreover, contributions to projects are more likely to be accepted (i.e. pull requests) if the contributor and evaluator are from the same country (Rastogi et al. (2018)). In general contributions for less developed countries, besides being fewer in number, are more likely to be rejected (Furtado et al. (2020)).

In some sense it is remarkable that geography remains so important in OSS given that several of the mechanisms that explain why firms cluster do not clearly transfer to this specific context. For example, users and competitors in the OSS space are likely more geographically dispersed than in traditional industries. It seems that other mechanisms behind cluster strength and persistence, for instance network effects and knowledge spillovers, are enough to keep activity clustered even within the virtual, global collaboration network of OSS developers (Fershtman and Gandal (2011)). Contributors often only work on OSS in their free-time, next to a primary occupation, often in the software industry. In this way, the geographic concentration of OSS extends the concentration present in the software industry. The continued dominance of specific clusters, once they emerge, may be less of a puzzle: developers working in such places have significant advantages over their counterparts in peripheral regions. Indeed norms and practices of the software industry tend to flow from core to periphery locations (Takhteyev (2012)). At the same time, projects on GitHub have a variety of geographic collaboration patterns, which are constantly evolving

(Heller et al. (2011)).

Local Effects of OSS Development

We have suggested that OSS development, like any other frontier knowledge creation, occurs in geographic clusters. At the same time OSS activity is generally unpaid and its outputs are public goods, distributed world wide via the Web. And though OSS contributes substantial amounts to global GDP and productivity in this way (Greenstein and Nagle (2014)), it begs the question whether locations hosting OSS clusters benefit in the same way as they might from hosting, for example, a cluster of highly profitable biotech companies. While OSS activity does not directly generate taxable profits, recent work has shown that it generates significant benefits that accrue locally (Nagle (2019b); Nagle et al. (2020)).

Specifically, the public nature of OSS activity fosters entrepreneurship and investment by providing strong signals of quality (Wright et al. (2020)). It also creates opportunities for developers and firms to learn by doing and to better integrate user feedback. Because of its open nature, OSS developers can effectively signal their specific areas of expertise, broadening the network of potential collaborators in a new venture. Public and transparent information about how people and teams work together makes it easier to attract outside investments (Kaminski et al. (2019)), to get quality feedback (Wachs and Vedres (2021)), and to learn by observing (Riedl and Seidel (2018)). A large OSS footprint also suggests that a location has complementary assets necessary for software entrepreneurship, which can contribute to network effects. This presents OSS activity as an effective proxy measurement for a place's technological development and capacity for IT innovation.

Beyond signaling effects, participation in OSS also brings many benefits. Firms contributing to OSS capture more value from the resulting software than free-riding competitors because they learn from the experience and gain valuable feedback (Nagle (2018)). In general OSS is high quality because it benefits from the attention of a large, relatively independent crowd of users (Aksulu and Wade (2010); Lakhani and Von Hippel (2004); Raymond (1999)). Firms who use OSS benefit from this ecosystem (Nagle (2019b)). Indeed these factors have led to growing adoption of and participation in OSS, including by firms with a traditionally proprietary orientation such as Microsoft and Apple (Anthes (2016)). Companies increasingly release in-house projects as OSS, for example Google's TensorFlow library for machine learning and Facebook's React web framework. Though companies cannot prevent competitors from using work that their employees contribute to OSS projects, they encourage such contributions anyway to gain legitimacy and access to communities (Dahlander and Wallin (2006)). OSS development can also serve as an incubator for new software ventures, recalling the "doing-using-interacting" model of innovation (Alhusen et al. (2021); Jensen et al. (2007)). These virtuous effects of OSS activity manifest locally, through productivity growth in firms and new ventures.

Besides the specific outputs of the OSS sector and their direct uses, software know-how itself is clearly an important input in emerging sectors including Industry 4.0 (Balland and Boschma (2021)) and artificial intelligence, for instance via the intensely computational methods of deep-learning (Klinger et al. (2021)). It also contributes to productivity gains in manufacturing (Branstetter et al. (2019)), suggesting that software has a complementary role in many sectors of the modern economy (Neffke (2019)), including for example biotechnology via bioinformatics (Hu et al. (2021)). In this way local activity in OSS, which is more easily observable than closed source software development, can serve as a valuable signal of the potential of a city, region, or country in the digital economy.

Public Policy and OSS

Given its economic impact, it is not surprising that policymakers are

increasingly interested in promoting OSS activity. These efforts often draw on insights from the cluster policy literature on traditional industries, but it is unclear if these ideas transfer to the OSS context. Many OSS contributions come from independent individuals, working in their spare time. OSS contributors only rarely receive direct financial support for their work, for instance through crowdfunding (Overney et al. (2020)) or consulting (Eghbal (2020)). Indeed they tend to be motivated by social and reputational reasons, and by specific technical problems (Gerosa et al. (2021)). Economic gains from OSS activity tend to accrue only in the long run (Lerner and Tirole (2002)) and so are rarely counted on by contributors. So while cluster policy mainstays like SME tax incentives, networking events, and support for fundraising and branding (Uyarra and Ramlogan (2016)) may help crystallize the economic potential of an active OSS scene, they are unlikely to attract or create new OSS contributors. In this way, optimal cluster policy for OSS activity merits additional investigation.

On the other hand, there has been significant effort by governments to encourage OSS activity in general. Lee (Lee (2006)) and Blind (Blind et al. (2021)) group motivations for public sector support of OSS into four groups: economic, technological, legal, and political. For example, using OSS saves on direct software costs and avoid vendor lock-in. Technologically, OSS solutions may be superior to closed source alternatives. Legally, governments, especially in the developing world, can avoid issues around software piracy by using OSS solutions. Finally, OSS use has political reasons: it increases transparency and eases public access to government IT infrastructure. These direct motivations and others have led to a significant amount of policy support OSS use.

An example of a successful policy promoting the use of OSS is the French law Circulaire 5608, passed in 2012 (Blind et al. (2021)). Motivated by potential cost-savings of avoiding closed source software licensing fees, the law requires French public bodies to consider open source solutions when procuring software. A study of the impact of this law by Nagle found significant social welfare gains (Nagle (2019a)) on the order of tens of millions of Euros. More significantly, second order effects observed include an increase in IT start-ups (9-18% yearly increase vs. the counterfactual) and employment in IT employment (7-14%), and a decrease in IT patents (5-16%). In general, policy promoting OSS has focused on government purchasing and public sector use of software while less has been done to incentivize private sector or individual contributions to OSS (Blind et al. (2021)), let alone in terms of cluster policy.

So while governments tend to have policy on OSS, it tends to focus on how the public sector can cut costs and improve transparency by using OSS, rather than how OSS adoption and development can improve local productivity and innovation outcomes. Public policy on the use of OSS also tends to happen at the national level, even though we will show that activity tends to cluster in regions. The rest of this work presents data on the actual geographic distribution of OSS developers and the characteristics of those places, both national and regional, which have many developers. Reflecting on these geographic distributions, we will revisit the question of effective policies for the promotion of OSS development in our discussion.

Research Design and Data

We now describe our data collection and processing pipeline. Before gathering information about activity levels and positions of individual OSS contributors, one needs to define contributors and their activity. Developers share code using version control protocols – allowing them to track changes, compare, test and merge with modifications of others – on specific online platforms. The most widely used protocol is called *git*, and the most widely used public platform for projects using *git* is *GitHub*. Modifications to files are collected in *commits*, which can be seen as snapshots of code edits. We use *commits*, as elemental contributions in OSS, to quantify the level of activity of individual developers.

Data on commits contributed to public projects on GitHub is made

Table 1

Country shares of active OSS contributors on GitHub in 2021. We include the top 30 countries and compare our data with similar snapshots reported in previous work using Sourceforge (2008) [Gonzalez-Barahona et al. \(2008\)](#) and GitHub (2010, only top 10 available) [Takhetev and Hiltz \(2010\)](#). Across countries the distribution has become more uniform. South and East Asian and Latin American countries have seen the greatest relative increase in share of global OSS contributors.

Rank	Sourceforge 2008		GitHub 2010		GitHub 2021		Rank Chg. vs. 2008
	Country	Share	Country	Share	Country	Share	
1	United States	36.1	United States	38.7	United States	24.6	-
2	Germany	8.1	UK	7.7	China	5.8	+4
3	UK	5.1	Germany	6.2	Germany	5.6	-1
4	Canada	4.2	Canada	4.3	India	5.4	+7
5	France	3.8	Japan	3.9	UK	5.0	-2
6	China	3.1	Brazil	3.6	Brazil	4.4	+6
7	Australia	2.7	France	3.2	Russia	4.3	+6
8	Italy	2.6	Australia	3.1	France	3.8	-3
9	Netherlands	2.5	Russia	2.3	Canada	3.8	-5
10	Sweden	2.0	Sweden	2.2	Japan	2.7	+5
11	India	1.9			South Korea	1.9	+14
12	Brazil	1.8			Netherlands	1.8	-3
13	Russia	1.6			Spain	1.8	+1
14	Spain	1.6			Poland	1.8	+2
15	Japan	1.3			Australia	1.8	-8
16	Poland	1.2			Sweden	1.2	-6
17	Belgium	1.2			Italy	1.2	-9
18	Switzerland	1.0			Ukraine	1.2	New
19	Austria	0.8			Switzerland	1.2	-1
20	Denmark	0.8			Indonesia	1.0	New
21	Singapore	0.8			Taiwan	0.8	+9
22	Finland	0.8			Colombia	0.8	New
23	Norway	0.7			Argentina	0.7	+4
24	Mexico	0.7			Mexico	0.7	-
25	South Korea	0.7			Norway	0.7	-2
26	Israel	0.6			Belgium	0.7	-9
27	Argentina	0.6			Denmark	0.7	-7
28	Hungary	0.6			Finland	0.6	-6
29	Romania	0.5			Vietnam	0.6	New
30	Taiwan	0.5			Austria	0.6	-11

available and dynamically updated on the GH Archive database¹. We use this database rather than querying data from the GitHub API. The next step is to assign GitHub accounts of authors to their commits. Commits themselves contain plaintext names and email addresses of authors, which do not correspond directly to GitHub accounts. For instance, a GitHub account user may contribute commits from multiple computers, each linked to git via a different name and email. Merging these identities under one developer is a crucial part in geolocating GitHub users, as the clearest information about their geographic location is provided on their GitHub profile page. We therefore applied a method from the software engineering research community to link email addresses and specific commits to GitHub user accounts (which we assume correspond to individual developers) ([Montandon et al. \(2019\)](#)) using the GitHub API: for each email address, we select a random commit made by that email address and query the GitHub API to retrieve the specific account login associated to the commit. In case the API cannot resolve the account using this first commit, we try three additional commits submitted using this email. As we are interested in active GitHub contributors, we consider all email addresses with at least 100 commits over the two years of 2019 and 2020. This corresponds to an average of nearly one commit per week. We note that our subsequent results replicate completely when applying a stricter threshold of at least 200 commits for inclusion. The results of this analysis are available upon request.

Having associated GitHub accounts to contributions, we access information about users via the GitHub API. In particular we access the *location* and *Twitter account* of individual users, when provided. Through the Twitter API, one can also retrieve location information of Twitter users, when provided. A third way of gathering information on location is through email suffixes, belonging either to a country or institutions such as universities. We describe our method of geolocating developers

in the following section.

Geolocation

Given a collection of active GitHub contributors and their account information, our goal is to infer the location for as many users as possible. We first focused on the raw location fields provided on GitHub user profiles. We selected the Bing Maps API to carry out our geocoding task. This API resolves multiple input languages (“Vienna” and “Wien” refer to the same location) and can handle inputs at varying scales from country to geolocation precise to within meters. It also appropriately handles edge cases of geocoding online profiles, for instance that users may give unreal or sarcastic locations (“the moon”) ([Hecht et al. \(2011\)](#)).

In case a user did not share a location on GitHub, or the Bing API was not able to geolocate the string that the user did share, we check whether the user linked to a Twitter account. If so, we attempted to geolocate the Twitter account in a similar way, using the location field provided by the user on Twitter. In case we could not geolocate a user from their Twitter data, we considered the email suffixes under which they made commits. Email suffixes can suggest the location of individuals in two ways: by the country domain (i.e. .jp for Japan, .it for Italy) or by a university suffix. In the latter case we imitate previous work associating GitHub contributors to universities ([Valiev et al. \(2018\)](#)) using a list of universities, their locations and email domains maintained by Hipo². We only infer user country from email suffix data.

Using this pipeline we could geolocate 587,852 active OSS contributors (out of 1,124,874 accounts with at least 100 commits) to at least the country level. 502,415 or 85% user locations were identified from their GitHub account information alone. In other words, by considering

¹ <https://www.gharchive.org/>

² <https://github.com/Hipo/university-domains-list>

Table 2
Countries ranked by number of OSS developers per capita, top 50.

Rank	Country	ISO2	Count Total Contributors	Pop. (mm)	Cont. / 100k
1	Iceland	IS	421	0.4	105
2	Switzerland	CH	7197	8.6	84
3	Norway	NO	4012	5.3	76
4	Sweden	SE	7323	10.3	71
5	Finland	FI	3813	5.5	69
6	Denmark	DK	3906	5.8	67
7	Netherlands	NL	10773	17.3	62
8	Canada	CA	22269	37.6	59
9	Estonia	EE	760	1.3	58
10	Luxembourg	LU	324	0.6	54
11	New Zealand	NZ	2642	4.9	54
12	Singapore	SG	3102	5.7	54
13	Ireland	IE	2531	4.9	52
14	United States	US	144371	328.2	44
15	United Kingdom	GB	29452	66.8	44
16	Australia	AU	10337	25.4	41
17	Germany	DE	33212	83.1	40
18	Austria	AT	3276	8.9	37
19	France	FR	22551	67.1	34
20	Belgium	BE	3935	11.5	34
21	Israel	IL	2488	9.1	27
22	Belarus	BY	2532	9.5	27
23	Portugal	PT	2802	10.3	27
24	Lithuania	LT	748	2.8	27
25	Poland	PL	10406	38.0	27
26	Czechia	CZ	2805	10.7	26
27	Bulgaria	BG	1755	7.0	25
28	Slovenia	SI	492	2.1	23
29	Latvia	LV	443	1.9	23
30	Spain	ES	10593	47.1	22
31	Malta	MT	112	0.5	22
32	Taiwan	TW	4979	23.6	21
33	South Korea	KR	10921	51.7	21
34	Hungary	HU	1813	9.8	18
35	Croatia	HR	742	4.1	18
36	Russia	RU	25271	144.4	18
37	Hong Kong	HK	1303	7.5	17
38	Ukraine	UA	7204	44.4	16
39	Serbia	RS	1039	6.9	15
40	Cyprus	CY	168	1.2	14
41	Greece	GR	1510	10.7	14
42	Slovakia	SK	719	5.5	13
43	Japan	JP	15706	126.3	12
44	Uruguay	UY	435	3.5	12
45	Brazil	BR	25891	211.0	12
46	Costa Rica	CR	593	5.0	12
47	Italy	IT	7204	60.3	12
48	Romania	RO	1979	19.4	10
49	Namibia	NA	260	2.5	10
50	Argentina	AR	4332	44.9	10

email suffixes and Twitter data we could increase our pool of geolocated developers by 15%. As country-identifying email domains and Twitter use vary significantly between countries, we share data on the number of users identified by each of the three methods in the iterative process. We note that a share of users (30%) could only be geocoded at the country level, for instance those classified using email-suffix data or giving only coarse geographic information in their location fields (i.e. “Austria”). Specifically, we could infer subnational locations for 415,783 users. Code to replicate our data mapping pipeline is available at: <https://github.com/n1tecki/Geography-of-Open-Source-Software>.

Data Availability

A primary goal of this work is to make geographic data on OSS developers accessible to researchers. We have uploaded both national and regional datasets to GitHub, available at: https://github.com/johanneswachs/OSS_Geography_Data. In order to protect user privacy, we only share geographically aggregated counts of active users. For example, one file includes the number of active developers located in each of the

50 US States in 2021. In particular we share data in CSV files comparing counts of active OSS contributors across Countries, European NUTS2 regions, and sub-national units (states/provinces) of the US, Japan, China, India, Russia, and Brazil.

Analysis and Results

We begin by reporting summary statistics on the number of developers we located in various countries and regions. We highlight which countries host the most OSS developers both in raw terms and per capita. We compare our results with those from previous works published in 2008 and 2010, studying the geographic distribution of developers on the Sourceforge and GitHub platforms, respectively. We also present evidence that the count of active OSS developers correlates strongly with a variety of measures of development and quality of living indicators, above and beyond economic development. We replicate these findings at the regional scale using data from the European NUTS2 regions. We then present an analysis of the geographic concentration of OSS developers within countries, finding that they are in general highly concentrated.

International Comparison

In Table 1 we report the top thirty countries, ranked by overall share of active OSS developers, and compare data from 2021 with snapshot data from previous work, carried out over 10 years ago. At first glance, our results indicate significant changes in the global distribution of OSS activity since 2010 (Gonzalez-Barahona et al. (2008); Takhteyev and Hilt (2010)). While North American and Western European countries are still leading locations for OSS, Asian, Eastern European, and Latin American countries are catching up. In Table 2 we report the top 50 countries by active OSS developers per 100k inhabitants. Here the top ranks are dominated by small and wealthy European countries.

A more interesting ranking of national activity in OSS would take into account both the population of each country and its level of economic development. In Figure 1 we present the relationship between income per capita, sourced from the World Bank, and the number of active OSS developers per million inhabitants, both on logarithmic scales. We exclude countries with a population of less than one million people for the sake of visualization. The regression fit explains roughly two-thirds of the variance among all countries, but only 40% for countries with an income per capita of at least \$10,000. Countries above the regression line have more OSS developers per capita than expected for their level of economic development, while those below have less. Ukraine, Belarus, Namibia, Brazil, Bulgaria, and Estonia have more OSS activity than expected, while oil-rich states like Qatar, Kuwait, and Saudi Arabia are OSS laggards.

What explains these residuals? The ability and decision to contribute to OSS projects is likely a complex and multifaceted process (Gerosa et al. (2021)), but we can compare the relative importance of various structural factors in a regression framework. Beyond the broad economic development of a country, measured by income per capita, internet penetration (The International Telecommunication Union (2019)) likely plays an important role. The UN’s Human Development Index (HDI) offers a broad measure of social development, including access to education and health services - likely important upstream factors facilitating OSS contributions (HDI (2019)).

One distinguishing aspect of OSS, compared to many other kinds of knowledge-intensive activities, is that OSS outputs are essentially public goods - goods that can be used by anyone. Though a thorough review of the economics of OSS (Lerner and Tirole (2002); Sahay (2019)) and individual motivations for contributing (Gerosa et al. (2021)) is beyond the scope of this article, we do expect that OSS activity will be higher where people are more inclined to contribute to public goods. For instance, people living in areas with high levels of generalized trust, that is to say where individuals are more likely to report that in general “most

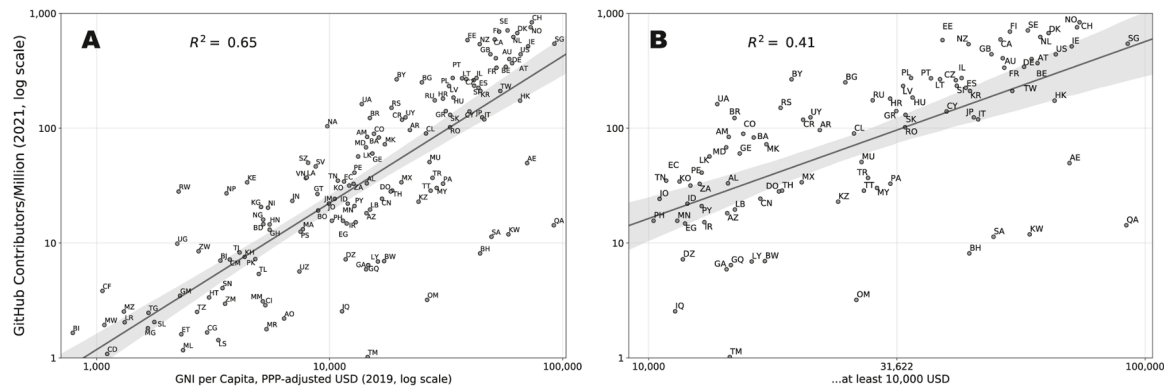


Fig. 1. Country economic development (PPP-adjusted GNI per capita) and OSS contributors per capita, log-log scale. A) We observe a strong relationship between economic development and OSS activity, though some countries deviate significantly from the trend. B) Zooming in on wealthier countries only, the relationship weakens significantly, suggesting other factors play an important role in OSS activity. For sake of visualization, we exclude countries with fewer than 1 million inhabitants or less than 1 contributor per million inhabitants.

Table 3

Spearman rank correlations between country-level social and economic development indicators and active OSS contributors per capita in 2021.

Country Feature	Spearman ρ	p-Value	Observations
PPP-adjusted GNI per capita (2019)	0.79	< 0.01	176
Human Development Index (HDI, 2019)	0.86	< 0.01	178
WVS share “most people can be trusted”	0.68	< 0.01	75
Index of Public Integrity (IPI, 2019)	0.87	< 0.01	117
Economic Complexity Index (ECI, 2019)	0.82	< 0.01	175
Deep-learning/AI publications/capita	0.67	< 0.01	183
Internet Penetration (2019)	0.78	< 0.01	180

Table 4

Regression models (1-5) relating country-level counts of GitHub contributors per million inhabitants (log-transformed) and socio-economic indicators. While income and internet penetration alone account for nearly two-thirds of variance in OSS activity (1), human development (2), quality of political institutions (3), and economic complexity (4) significantly improve model fit above and beyond that baseline. A combined model (5) explains over 80% of variance. We report robust standard errors.

	Active GitHub Contributors per Million Inhab. (log, 2021)				
	(1)	(2)	(3)	(4)	(5)
PPP GNI per Cap. (‘000 USD, 2019)	0.017*	-0.007	-0.002	0.004	-0.010*
	(0.009)	(0.007)	(0.009)	(0.008)	(0.006)
Internet Penetration (% of Pop., 2019)	0.043***	0.002	0.029***	0.026***	-0.003
	(0.005)	(0.008)	(0.005)	(0.006)	(0.009)
Population (log, 2019)	-0.145***	-0.071	-0.016	-0.143**	-0.038
	(0.052)	(0.044)	(0.046)	(0.056)	(0.057)
Human Development Index (2019)		11.709***			9.327***
		(1.528)			(1.553)
Index of Public Integrity (2019)			0.704***		
			(0.127)		
Economic Complexity Index (2019)				0.962***	0.683***
				(0.184)	(0.153)
Observations	174	173	115	150	149
Adjusted R ²	0.631	0.747	0.819	0.740	0.804
Residual Std. Error	1.191	0.986	0.788	1.016	0.882
F Statistic	81.3***	175.2***	195.8***	142.4***	155.6***

*p<0.1; **p<0.05; ***p<0.01

people can be trusted”, are known to be more likely to contribute to public goods (Rothstein and Uslaner (2005)). We use data from the most recent wave of the World Values Survey to measure this concept of generalized trust (Haerper et al. (2017)). Another feature of countries that strongly correlates with individual propensity to contribute to public goods is quality of government. We therefore also relate OSS outcomes to the Index of Public Integrity (IPI), a measure of the quality of public institutions in a country (Mungiu-Pippidi and Dadaşov (2016)).

Lastly, we consider the overall economic focus and specialization of a country. As noted before, oil-producing states seem to have fewer OSS contributors relative to their economic development. The extent to which countries specialize in more complex, coordination-intensive industries likely correlates significantly with OSS activity. To capture this aspect, we consider the Economic Complexity Index (ECI) (Hidalgo (2021)), a measure of the sophistication of a country’s export profile. To measure national sophistication in a specific frontier and software-adjacent field, we use a count the number of academic research preprints on AI/deep-learning published by each country in the last decade (Klinger et al. (2021)).

Each of these social, political and economic features has a strong correlation with the local intensity of OSS contributors. We report the Spearman rank correlation coefficient (ρ) of each variable with OSS contributors per capita in Table 3. We report a full correlation matrix and basic summary statistics in the appendix. Though it is not possible to disentangle the cause and effect relationships between these variables using observational data, we can observe how some of these variables mediate each other, and how they can explain a greater share of inter-country variance in OSS contributors. To do so, we regress the (log-transformed) number of OSS contributors per million inhabitants on a selection of these variables, reporting the results of several alternative specifications fit using ordinary-least-squares (OLS) in Table 4, reporting robust standard errors and measuring the goodness of model fit using adjusted R-squared.

Our baseline model, including per capita income, internet penetration, and population, explains around two-thirds of variance as measured by adjusted R-squared. Adding HDI, IPI, or ECI as individual features increases the variance explained by over 10%. All three variables have a significant positive relationship with great OSS activity in a country. A model combining the baseline model with HDI and ECI accounts for over 80% of variance in OSS activity. In general we observe both a statistical significant relationship between these socio-economic factors in our models, and a large increase in the quality of the model fit when including them. Given that our features have significant pairwise correlations, we test for issues of multicollinearity which may influence our estimates using Variance Inflation Factor (VIF) tests for each specification. The results, reported in the appendix, indicate moderate levels

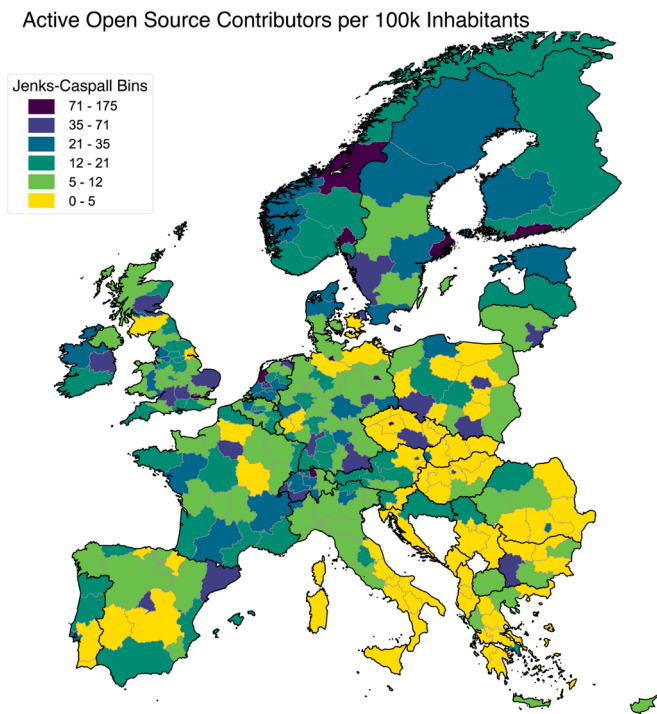


Fig. 2. Active OSS developers concentrations in early 2021 per 100,000 inhabitants, NUTS2 regions. We observe significant within country variation.

of correlation within standard bounds.

These models suggest that OSS activity is not merely the by-product of an advanced economy, but also depends to a significant degree on social, educational, and political institutions and the degree of technological sophistication in a country. These results shed light on the residuals of the simple bivariate relationship between economic development and OSS activity reported in Figure 1. To take an extreme example, though Estonia and Bahrain have comparable average income, the density of OSS contributors in Estonia is nearly two orders of magnitude greater than that of Bahrain. Much of this difference is captured by variation in, for example human development or economic complexity. But even Latvia and Lithuania, geographically and historically linked to Estonia, have significantly fewer OSS developers - suggesting rich nuance in the variation of OSS activity and scope for specific policy interventions, to be discussed later in the paper. First, however, we will zoom in on the sub-national level. As we will see in the following section, geographic variance in OSS activity becomes more difficult to explain with macro indicators at finer spatial scales.

Regional Variation

Comparing OSS activity between countries indicates that it has spread internationally to a significant extent in the last ten years. However, we know little about the distribution of activity within countries. As mentioned above, while most prior work has focused on international comparisons, Takhteyev and Hilts (Takhteyev and Hilts (2010)) reported data on local clusters in their work from 2010. They estimated that 7.4% of global contributors were in the San Francisco Bay Area, suggesting an immense local concentration of OSS activity. In this section we explore the local distribution of OSS developers in various countries. We focus first on European NUTS2 regions. We again relate socio-economic features to OSS activity, replicating our international findings at the regional scale. We also report data on the top US metropolitan statistical areas. We then consider the *concentration* of developers within multiple countries including the EU, US, China, India, Japan, and Brazil. We find that OSS activity is significantly concentrated

relative to the distribution of the general population in all countries we examine, though with significant heterogeneity. Repeating the calculation using university educated workers or workers employed in high-tech. fields instead of OSS contributors, we find far lower levels of concentration.

European Regions

We zoom in on European NUTS2 regions. These regions are especially useful because we can compare regions from multiple countries with generally consistent statistics, sourced from Eurostat. In particular we use the regions defined in 2016. In Figure 2 we map the number of OSS developers per 100,000 inhabitants by European NUTS2 region using Jenks-Caspall bins. We note that the five London NUTS2 regions are merged into a single unit because developers tended to refer to their location as London rather than “Inner London”. We can observe several patterns. First we note that major hubs such as London, Amsterdam, Berlin, Prague, Zurich, Hamburg, Helsinki, Oslo and Stockholm tend to have significant OSS presence. Second, we can see significant variation between regions within countries. Some countries have regions in both the lowest and highest valued bins. In other countries, such as Italy, Spain, and France, the distribution seems to be more uniform. We will return to an investigation of the within-country spatial concentration of developers later.

Knowing the locations of OSS developers at the regional levels, we can attempt to replicate our earlier findings about the relationship between socio-economic development indicators and OSS activity. Again the goal is to show that OSS activity is related to more than just economic development outcomes, albeit this time at a sub-national scale. On top of a baseline set of features including internet penetration, GDP per capita, population and density, we consider the relationships between various indicators of social and technological development and OSS activity. In particular we consider general levels of social trust, measured by the European Values Survey (GESIS Data Archive (2017)), R&D Spending per capita, share of workers employed in high-tech industries, number of patents per 100k inhabitants (sourced from the OECD REGPAT database (Maraut et al. (2008))), and share of the prime working-age population with tertiary education. Unless otherwise noted data are sourced from the 2017 QoG Basic dataset (Dahlberg et al. (2021)). When data on any feature is only available at the coarser NUTS0 (country) level, we impute from the country level to the regions. Finally we also consider the number of patents filed in the Electrical Engineering (EE) sector in case the relationship between OSS activity and closed-source innovation is heterogeneous across sectors. The WIPO IPC Technology Concordance Tables categorize patents into five fields at its coarsest level, of which EE is most directly related to software³.

We would like to model the relationship between these indicators of socio-economic development and OSS activity per capita at the NUTS2 level, observing both statistical significance and the share of variance explained. We again constructed a baseline model (including internet penetration, GDP per capita, population, and population density) and added the other features one at a time. However, when modeling outcomes of geographic units at finer spatial scales, it is important to consider spatial correlations in variables of interest (Cliff and Ord (1981)). Such correlations violate OLS assumptions about the independence of the error term and can introduce bias. To test for spatial autocorrelation in our data and modeling set up, we first calculated Moran's I on the distribution of OSS developers per 100k inhabitants (logged) across the NUTS2 regions. Two regions are considered adjacent if they share any border, including corners. We found evidence of significant spatial correlation ($I = 0.42, p < .001$).

This suggests that linear regressions fit by OLS may suffer from bias

³ The other categories are: Instruments, Chemistry, Mechanical Engineering, and Other.

Table 5

GMM spatial regression models (1-7) (Arraiz et al. (2010)) relating EU NUTS2 counts of GitHub contributors per 100k inhabitants (log-transformed) and socio-economic indicators. Income, population, and internet penetration account for just over one third of variance in OSS activity (1). Social trust (2), R&D spending (3), employment in high tech sectors (4), innovation activity (5,6), and higher education (7) all explain additional variance above this baseline (from 1 to 14%). In each model the spatial autoregressive term lambda is positive and significant, indicating a positive adjacency relationship: neighboring regions tend to have similar levels of OSS activity even accounting for the features in each model.

	Active GitHub Contributors/100k Inhab. European NUTS2 (log, 2021)						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Internet Penetration	0.010**	0.005	0.012***	0.012***	0.012**	0.011**	0.004
—(% of Pop. 2017)	(0.004)	(0.005)	(0.004)	(0.004)	(0.004)	(0.004)	(0.004)
GDP per Cap.	0.600***	0.609***	0.012	0.367**	0.248	0.269	0.352**
—(log Eur, 2017)	(0.140)	(0.154)	(0.256)	(0.147)	(0.176)	(0.202)	(0.151)
Population	0.143*	0.178	0.356***	0.166*	0.063	0.089	0.207**
—(log, 2020)	(0.084)	(0.113)	(0.131)	(0.092)	(0.086)	(0.085)	(0.081)
Population Dens.	0.043*	0.036	0.044	-0.018	0.055**	0.050*	0.030
—(log, 2017)	(0.023)	(0.024)	(0.029)	(0.022)	(0.027)	(0.028)	(0.022)
EVS Trust		0.315**					
—(2017)		(0.151)					
R&D Spend. per Cap.			0.113**				
—(log, 2017)			(0.051)				
% Empl. High-Tech				0.083***			
—(2019/20)				(0.011)			
Patents Elec-Eng./100k					0.058***		
—(log, 2017)					(0.022)		
Patents/100k						0.047*	
—(log, 2017)						(0.026)	
% with Tertiary Edu.							0.021***
—(2019/20)							(0.002)
Lambda	0.099***	0.114***	0.073***	0.083***	0.078***	0.084***	0.071***
—(est. spatial dep.)	(0.014)	(0.012)	(0.015)	(0.015)	(0.013)	(0.014)	(0.014)
Observations	276	198	258	262	258	258	276
Pseudo-R ²	0.392	0.429	0.417	0.509	0.411	0.399	0.536

*p<0.1; **p<0.05; ***p<0.01

introduced by spatial autocorrelations. To test this, we estimated these models and applied diagnostic tests (Moran's *I* of the residuals and Lagrange Multiplier tests (Anselin (1988))). These tests, reported in the Appendix, confirmed that our estimates are significantly biased by spatial autocorrelation (i.e. the p-value of Moran's *I* calculated on the residuals was always below 0.1, and below 0.05 for a majority of the models). We therefore used a general method of moments (GMM) modeling approach developed by Arraiz et al. that corrects this bias and adjusts for spatial correlation in both the dependent variable and independent variables (Arraiz et al. (2010)). We used the Pysal implementation of this model in the Python programming language (Rey and Anselin (2010)).

We report the results of our GMM regression analysis in Table 5 with the (log-transformed) number of active GitHub contributors in a region per 100k inhabitants as the dependent variable. The baseline model again indicates that economic development and internet access have a significant positive relationship with OSS activity in regions. However, at this spatial scale the model has a significantly less accurate fit (pseudo $R^2 \approx .39$) than a similar model predicting OSS activity at the country level. This observation applies also to the feature-rich models. Though generalized trust, R&D spending per capita, share of employment in high-tech industries, share of population with tertiary education, and patents are significant predictors of greater OSS activity, the overall model fit only improves somewhat when adding these features. Only when we include share of the working age population with tertiary education or share of workers employed in high-tech sectors does the variance explained exceed 50%. While factors like the presence of technologically advanced industries and an educated workforce clearly relate to OSS activity, it seems that at the regional level, more idiosyncratic forces determine local participation in OSS.

The relationships between the two patenting variables and OSS activity also merit comment. Patents are awarded to protect intellectual

property and to block the uncompensated use of a creator's ideas. In a naive sense, patenting would appear to be a substitute for activity in open source and the creation of public goods. In practice, however, we see that patenting in electrical engineering, and to lesser extent patenting in all fields, has a significant positive relationship with OSS activity in regions. If OSS activity were to crowd-out patenting, we would expect to see the opposite relationship. This suggests that OSS plays a complementary role in the innovation process, likely via knowledge spillovers discussed earlier. Hybrid outcomes are also possible, in which software that accompanies a proprietary product is made open source. More work is needed to understand the potential impact of policy interventions on the relationship between open source activity and patenting. This finding also suggests how OSS activity can serve as a proxy of useful skills cognitively close to those involved in closed-source innovation (Boschma (2005)).

We have seen that on a regional level, OSS activity is positively related with economic development, activity in technology intensive sectors, and the presence of an educated and trusting population. At the same time, models including these factors fail to explain over half of the observed variance in OSS activity between regions in Europe. This suggests that while a place can have the right ingredients for an OSS hotspot, for instance if it hosts a wealthy, well-educated, and digitally connected populace, and still fall short. These findings motivate our analysis of the concentration of OSS activity within countries, which will demonstrate that this variance is also large in size.

US Metropolitan Statistical Areas

Before turning to an analysis of within country concentration, we briefly present data on the distribution of developers in US metropolitan areas. Our reasons for this detour are twofold: first we can estimate the share of the global developer population in the San Francisco Bay Area

Table 6

Top 10 US Metropolitan Statistical Areas with at least one million inhabitants ranked by active GitHub developers per capita.

MSA Name	Count Contributors	Population	Contributors/100k
San Jose-Sunnyvale-Santa Clara, CA	4,587	1,990,660	230
San Francisco-Oakland-Hayward, CA	10,702	4,731,803	226
Seattle-Tacoma-Bellevue, WA	8,830	3,979,845	221
Austin-Round Rock, TX	3,370	2,227,083	151
Portland-Vancouver-Hillsboro, OR-WA	2,751	2,492,412	110
Boston-Cambridge-Newton, MA-NH	5,221	4,873,019	107
Denver-Aurora-Lakewood, CO	2,555	2,967,239	86
Raleigh, NC	1,159	1,390,785	83
Salt Lake City, UT	989	1,232,696	80
New York-Newark-Jersey City, NY-NJ-PA	11,579	19,216,182	60

and compare it with previous estimates from 2010. Second, the results indicate that concentration is also high in the US. While US states often have populations similar to European countries, and US cities often contain multiple counties, the US Census bureau defines Metropolitan Statistical Areas (MSAs) to highlight urban agglomerations.

In Table 6 we report the top 10 MSAs with at least one million inhabitants by the number of developers per capita. The two leading regions are San Francisco and Silicon Valley, while Seattle has a similar density. A back of the envelope calculation suggests that together the two Bay Area regions account for about 3.7% of all OSS developers world wide (summing 4,587 and 10,702, and dividing by 415,783, the number of OSS developers we could locate at a subnational level). This is

around half of an estimate of the share of OSS developers in Silicon Valley from 2010 of 7.8% (Takhteyev and Hiltz (2010)). We report a table of the top 50 MSAs by contributors per capita including those with smaller populations in the Appendix, noting that many university towns appear in prominent positions.

These results mirror recent work on the spatial distribution of patents in the US (Chattergoon and Kerr (2022)). In that work, Chattergoon and Kerr find that in the years 2015-2019, roughly one in three patents were made by individuals living in one of six metropolitan areas: San Francisco, Boston, San Diego, Seattle, Denver and Austin. Among software-related patents, their share was around 45%. Our data estimates that 34% of US-based open source developers live in one of these six cities. As in the European case, the geographic clustering of software developers closely tracks the most innovative regions of the US.

Concentration

As we have noted above, many kinds of knowledge-intensive activities are generally known to cluster in specific areas within countries. One of the primary goals of our study is to examine the extent to which this is true of OSS. Though it may be especially conducive to remote collaboration and decentralization, Figure 2 provides qualitative evidence that these aspects of OSS development do not outweigh the tendency of knowledge-intensive activities to cluster. We now introduce a measure to quantify this phenomenon, and compare the degree of geographic concentration of OSS activity in countries.

There are many measures of the dispersion or concentration of people or things across geographic regions (Ellison and Glaeser (1997)). As we are interested in comparing the relative concentration of OSS developers between countries, we need a measure which considers population heterogeneity between regions within countries. For example, in a hypothetical country with two regions, A and B containing 80% and 20% of the country's population, respectively, an 80%-20%

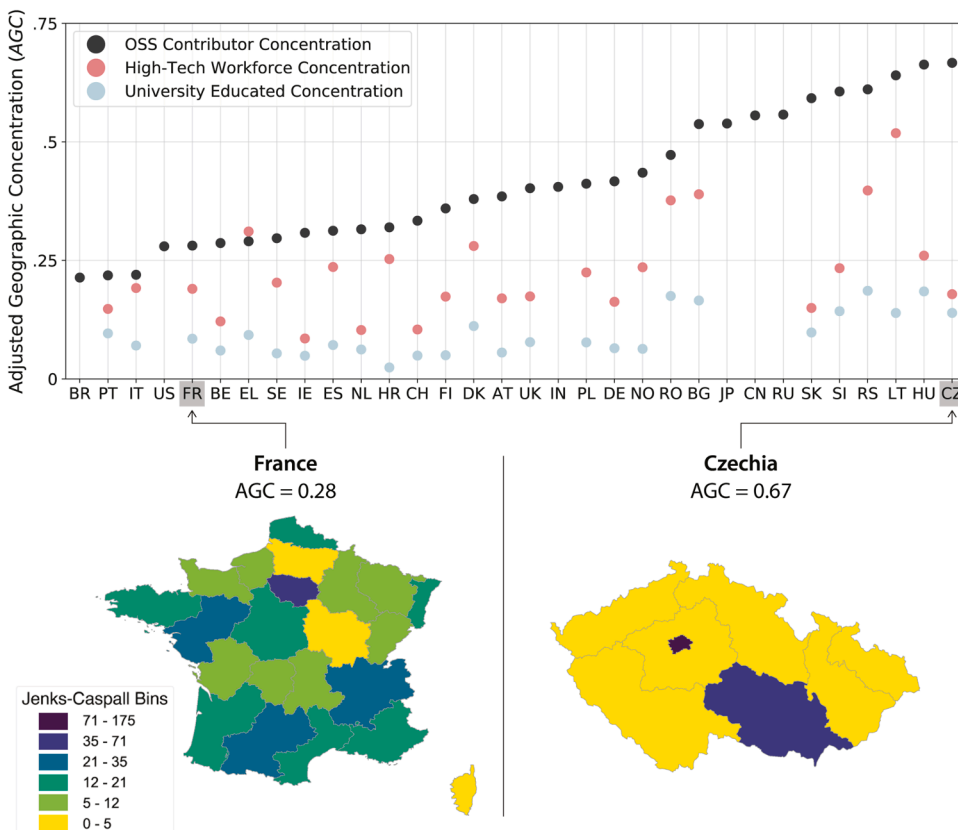


Fig. 3. OSS contributor regional concentration within countries, measured using Adjusted Geographic Concentration (AGC). A score of zero indicates that OSS contributors are distributed across a country's regions in proportion to population. AGC hypothetically equals 1 if all OSS contributors are concentrated in the least populated region of the country. For European countries with NUTS regions we compare the AGC of OSS contributors with the AGC calculated for university educated workers and people employed in high-tech sectors. Below: Examples of countries with relatively low (France) and high (Czechia) OSS AGC scores. Darker NUTS2 regions indicate more OSS contributors per capita, with bins as in Figure 2.

distribution of OSS developers between regions A and B should not be interpreted as concentration.

Measures like the Herfindahl Index depend on the number of regions, while the Ellison-Glaeser measure is sensitive to variance in population between regions (Ellison and Glaeser (1997)), and Gini-like measures quantify inequality, which is distinct from concentration. The *Adjusted Geographic Concentration* (AGC), developed by the OECD (Spiezia (2003)), measures concentration that is comparable between countries with different numbers of regions and different distributions of the underlying population between them (Rovolis and Tragaki (2006)).

Consider a country C with a population P split into N regions. The regions $i \in \{1, 2, \dots, N\}$ have shares of the population p_i . Denoting by m_i the share of OSS developers in a country living in the region G_i , we define the *Geographic Concentration* (GC) of developers in country C as:

$$GC(C) = \sum_{i \in C}^N |m_i - p_i|.$$

This measure sums the absolute differences in shares between the general population and the subpopulation of interest (in our case, active OSS contributors). This statistic tends to underestimate concentration in regions with a larger share of the population, and the validity of comparisons between countries with different numbers of regions is unclear. To address these issues the GC is usually scaled by its maximum possible value in each country: specifically, by the value it would take if all OSS developers were located in the least populated region of the country. In our previous notation:

$$GC_{max}(C) = (1 - p_{min}) + \sum_{i \in C, p_i \neq min}^N p_i = 2(1 - p_{min})$$

Dividing $GC(C)$ by $GC_{max}(C)$, we obtain the *Adjusted Geographic Concentration* (AGC) of a country C :

$$AGC(C) = \frac{GC(C)}{GC_{max}(C)}$$

The AGC varies between 0 and 1: a country in which the population of OSS developers is distributed in precisely the same proportions across region as the population, would have an AGC score of 0. A country in which all of OSS developers live in the region with the smallest population would have an AGC score of 1.

We calculated the AGC score for various countries, including European countries with at least 2 NUTS2 regions, and the US (states + DC, and MSAs), China (provinces and municipalities, excluding Hong Kong and Taiwan), India (states and union territories), Russia (federal subjects) and Brazil (federal states + the Federal District). We report our estimates of developer concentration by country in Figure 3.

We see that in all countries we examine, OSS development is concentrated regionally, relative to the general population. There is however significant variation between countries. For instance we can say that Brazilian, Portuguese and Italian OSS contributors are more evenly distributed amongst regions in those countries, than developers in Czechia, Hungary and Lithuania.

This analysis, however, does not make it clear how concentrated OSS developers are compared to other kinds of knowledge workers. In general regional statistics on such workers are not internationally comparable. However, among the European NUTS countries, we can make this comparison. We therefore recalculated the AGC of each country in this group, substituting the share of workers in high-tech sectors and with tertiary education, respectively, for OSS contributors in the calculation. If OSS contributors are more dispersed in a country than the university educated or high-tech workforce, we would expect the AGC to be higher under these alternative specifications.

In Figure 3 we observe the opposite effect: OSS contributors are significantly *more* concentrated in particular regions than either university educated or high-tech workers. This finding holds with

remarkable regularity across the countries we analyze (only Greece is an exception). We provide the full table in the appendix. These estimates of concentration also provide useful perspective for the previous more global analysis. The AGC scores for workers with higher education vary between .02 and .19 (mean: .09, stdev.: .05), for workers in high-tech sectors between .1 and .52 (mean: .23, stdev.: .10), and for OSS contributors in European countries between .22 and .67 (mean: .41, stdev.: .14). The AGC of OSS developers in all countries in our sample exceeds the average AGC of high-tech workers in European countries. Overall, these results present strong evidence that OSS developers cluster to a significant degree in all countries in our analysis.

While the spatial clustering of knowledge-intensive activity is in line with previous work, the observation that OSS development is generally more clustered in space than the high-tech workforce and the university educated is somewhat surprising. Contributions to OSS can be made from anywhere, are often on a volunteer basis, and interactions on GitHub between developers are often anonymous. Clusters of high-skilled workers usually tend to co-locate in neighborhoods, command high salaries, and have intensive networks (Florida (2002)). This suggests two potential causes for the clustering of OSS developers. The first is that perhaps OSS development is not as exceptional an activity as its surface-level characteristics suggest: for example OSS developers may derive (indirect) economic value from their activity (Lerner and Tirole (2002)). The alternative is that the social factors like peer influence, which depend to a significant extent on physical proximity (Latané et al. (1995)) and face to face contact (Storper and Venables (2004)), are what drive individuals to participate in OSS. Indeed recent work suggests that individual mentorship is one of the most effective ways to bring people into the OSS world (Steinmacher et al. (2021)). Likely both factors play some part in explaining the observed levels of concentration. We revisit the policy implications of these observations in the following section.

Conclusions and Implications

In this paper we presented an analysis of a novel dataset on the geography of open source software developers. We found that while the overall share of active developers has become more evenly distributed between countries, within-country regional differences remain strong. These heterogeneities are likely to persist: the social and economic spillovers of local OSS activity seem to be self-reinforcing. If indeed OSS is a driver of, and not merely a proxy for innovation outcomes, we need to better understand the role of actors such as universities, research institutes, and governments in promoting OSS activity (Nagle (2019a); Secundo et al. (2017)), for example in a mission-oriented context (Wanzenböck et al. (2020b)). While most policy supporting OSS activity focuses on the national level, our findings suggest that local and regional policy may be more appropriate.

Indeed, although digitalization facilitates collaborations across distances (Forman and van Zeebroeck (2019)), continued regional clustering suggests that location matters as much as ever. The network effects present in places such as Silicon Valley are strong enough to overcome other obstacles including higher tax rates and cost of living. In software knowledge spillovers and transfers still happen locally and within firms (Wu et al. (2018)). The developers in our dataset were geocoded in early 2021, roughly one year after the Covid-19 crisis went global. It remains to be seen whether proliferation of remote work and decentralization will be reflected in a change in the geographic distribution of OSS developers. If knowledge workers are going to permanently decamp to smaller cities in the post-Covid era, one would expect to see the first signs of this in OSS with its advanced infrastructure for remote collaboration. Our results, early in the post-Covid era, suggest that the winner-take-all dynamics of economic geography persist (Florida et al. (2021)).

Some limitations of our study highlight potential future work and extensions. The most obvious extension is to continue collecting data into the future to observe trends as they unfold, for instance to better

understand whether OSS activity follows or predicts entrepreneurial activity and innovation. Given the apparent interest in the future of work vis-à-vis Covid-19 and remote collaboration, this would be a valuable and important extension. More granular locations of developers, though difficult to source, would allow for studies of clustering of knowledge activity within cities (Hegyi et al. (2021)). The movement of individual developers could provide additional level of detail into how the tech world is adapting to OSS. Movements likely predict future inter-regional linkages, as those who arrive to a new place connect people in their new homes with those from their previous one (Kunczer et al. (2019)). Such data also presents the opportunity to study how software innovations diffuse geographically for instance via collaboration ties (Tóth et al. (2021)).

Our approach could be extended to cover alternative platforms for open source contributions including GitHub and Bitbucket, the absence of which may bias our results (Trujillo et al. (2021)). There are also potential biases within GitHub itself (Kalliamvakou et al. (2014)), for instance some projects are experiments, class projects, or webpages and not the kind of code that is widely used. Though we think this is unlikely to significantly bias our results, future work using our data should consider this potential limitation. Other possibilities for geolocating developers, including by geocoding the companies and organizations they work for, should be explored to add breadth. However, such extensions may bias results as developers in different countries may link to organizations at different rates. Nor does the use of an email address with a country top-level domain guarantee that an individual in question lives in that country. We have attempted to balance these concerns to provide an accurate measure of OSS activity at various scales.

Policy Implications

Perhaps the most important next step is to derive policy implications from our results. Though OSS activity is highly clustered now, that does not preclude other regions from increasing their participation in OSS and enjoying the apparent benefits of a strong OSS footprint. After all, regional clusters of creativity are known to rise and fall (Doehne and Rost (2021)) over long periods of time. We frame these implications along three lines, suggesting potential future work in each direction. First we consider what our results suggest about existing policies promoting OSS. We then consider which kinds of policies more generally, for instance those commonly discussed in the cluster policy literature, are likely to be effective. Finally, we discuss how OSS connects and relates to the rest of the economy.

As discussed in our literature review, national and regional public policy on the use of OSS tends to focus on the potential cost-savings of public sector adoption of OSS. As recent work has shown that local OSS activity has significant positive economic externalities, we suggest that there is ample scope for promoting OSS use in the private sector and at universities. For example, a regional scientific funding agency may encourage or require software outputs of projects to be released under an open source license. Students can be encouraged to contribute to OSS projects in the course of the projects as a way to build and demonstrate expertise. As a sector, OSS is in a unique position to span the branches of the so-called Triple Helix model of innovation policy that connects universities, private sector firms, and public sector institutions (Etzko-witz and Zhou (2017)). On a more fundamental level, our regression models suggest that human and social development has a strong relationship with OSS activity. From this finding we infer that it is necessary but not sufficient for a place to be wealthy and well-connected to the internet to develop a strong OSS community. Given that we can only explain roughly half of EU regional variance in OSS activity using socio-economic indicators, it seems that there are strong intangible drivers of OSS, suggesting that the right policies can make a difference.

The strong regional clustering observed in our data indicates that

insights from the cluster policy literature can inform policymaking on OSS. Effective cluster policy is often more about building networks and sharing information than providing specific economic incentives for an activity (Uyarra and Ramlogan (2016)). For example, a city may introduce a mentorship program linking experienced OSS developers with new developers; there is already evidence that one on one or small scale mentorship is one of the best ways to get new people involved in OSS (Steinmacher et al. (2021)). On a broader scale, cities and regions can support meetups of open source developers with space and resources. Fostering informal networks can make local ecosystems more resilient (Saxenian (1990)) and innovative (Trippel et al. (2009)). Firms can be informed about the various benefits of using and contributing to OSS (Nagle (2018)). On the other hand, policy efforts foster OSS use directly, for instance by mandates for public institutions, have had limited success (Blind et al. (2021)).

Startups and firms that create software could be informed about the potential gains of making their software open source (Nagle (2018, 2019b)). For example, a strong open source presence can be an effective signal of competence and ability to potential customers and investors. These kinds of policies represent a significant departure from the most common OSS policies implemented today. Between regions and cities there is also a potential to support networks of developers, not only projects, as the EU already does for traditional R&D sectors (Wanzenböck et al. (2020a)). Mapping these inter-regional networks of OSS developers would be a valuable next step.

We have also seen that OSS activity relates to innovation outcomes, namely patenting activity and the publication of research papers on AI. While we cannot demonstrate the direction of the effect in this case, i.e. whether OSS activity accelerates innovation activity or is merely a by-product of the process, the correlations we observe suggest this question merits additional study. One potential framework for such an investigation is the concept of relatedness, which suggests that places create new products and inventions by adapting and recombining existing expertise and know-how (Essletzbichler (2015); Hidalgo (2021)). From this perspective software expertise may open up specific and valuable new paths to innovation. Another related idea is the notion of technological complexity, which suggests that innovations that combine elementary ingredients in complex ways are crucial for economic growth at the knowledge frontier (Mewes and Broekel (2020); Pintar and Scherngell (2021)). In this context we suggest that software is a special kind of skill that enables people to combine inputs in novel ways (Gomez-Lievano and Patterson-Lomba (2021)). Understanding software's role in modern innovation, hence its impact on economic growth, is thus an area warranting additional research.

Declaration of competing interest

The authors declare that they have no conflict of interest.

Acknowledgements

JW, MN, and AP acknowledge support from the City of Vienna WU Jubilee Fund (Jubiläumsfonds der Stadt Wien). The authors thank Frank Neffke, Gergő Tóth, Sándor Juhász, and Balázs Lengyel for helpful discussions and Jeremias Brändle for assistance with the patent data.

Appendix

Summary Statistics and Robustness Tests

In this section we report numerous summary statistics about our data as well as robustness tests to supplement our regression analyses. This additional information is provided in Table 7, Table 8, Table 9, Table 10, Table 11, Table 12, Table 13, Table 14, Table 15, Table 16.

Table 7

Countries ranked by number of GitHub contributors (located via GitHub or Twitter location, or email suffix data), per 100k inhabitants. We exclude countries with fewer than 300k inhabitants and Montenegro, because the “.me” domain suffix is popular world-wide.

	Country	ISO2	# GitHub	# Twitter	# Email Suf.	Total Contributors.	Pop. (mm)	Contribs./100k
1	Iceland	IS	249	10	162	421	0.4	105
2	Switzerland	CH	4978	55	2164	7197	8.6	84
3	Norway	NO	3137	29	846	4012	5.3	76
4	Sweden	SE	6076	38	1209	7323	10.3	71
5	Finland	FI	3086	20	707	3813	5.5	69
6	Denmark	DK	2798	22	1086	3906	5.8	67
7	Netherlands	NL	8843	110	1820	10773	17.3	62
8	Canada	CA	19267	219	2783	22269	37.6	59
9	Estonia	EE	600	6	154	760	1.3	58
10	Luxembourg	LU	280	4	40	324	0.6	54
11	New Zealand	NZ	2299	28	315	2642	4.9	54
12	Singapore	SG	2818	33	251	3102	5.7	54
13	Ireland	IE	2224	25	282	2531	4.9	52
14	United States	US	128526	1831	14014	144371	328.2	44
15	United Kingdom	GB	24493	443	4516	29452	66.8	44
16	Australia	AU	8970	120	1247	10337	25.4	41
17	Germany	DE	25027	276	7909	33212	83.1	40
18	Austria	AT	2472	53	751	3276	8.9	37
19	France	FR	17474	202	4875	22551	67.1	34
20	Belgium	BE	3134	43	758	3935	11.5	34
21	Israel	IL	2131	43	314	2488	9.1	27
22	Belarus	BY	2375	8	149	2532	9.5	27
23	Portugal	PT	2485	32	285	2802	10.3	27
24	Lithuania	LT	683	5	60	748	2.8	27
25	Poland	PL	8864	51	1491	10406	38.0	27
26	Czechia	CZ	2771	34	0	2805	10.7	26
27	Bulgaria	BG	1510	11	234	1755	7.0	25
28	Slovenia	SI	411	6	75	492	2.1	23
29	Latvia	LV	371	4	68	443	1.9	23
30	Spain	ES	9091	157	1345	10593	47.1	22
31	Malta	MT	100	0	12	112	0.5	22
32	Taiwan	TW	4293	66	620	4979	23.6	21
33	South Korea	KR	10025	35	861	10921	51.7	21
34	Hungary	HU	1616	13	184	1813	9.8	18
35	Croatia	HR	666	3	73	742	4.1	18
36	Russia	RU	15543	108	9620	25271	144.4	18
37	Hong Kong	HK	1151	13	139	1303	7.5	17
38	Ukraine	UA	6941	29	234	7204	44.4	16
39	Serbia	RS	953	5	81	1039	6.9	15
40	Cyprus	CY	157	4	7	168	1.2	14
41	Greece	GR	1338	21	151	1510	10.7	14
42	Slovakia	SK	620	6	93	719	5.5	13
43	Japan	JP	12181	277	3248	15706	126.3	12
44	Uruguay	UY	397	11	27	435	3.5	12
45	Brazil	BR	24021	299	1571	25891	211.0	12
46	Costa Rica	CR	501	7	85	593	5.0	12
47	Italy	IT	5728	107	1369	7204	60.3	12
48	Romania	RO	1820	11	148	1979	19.4	10
49	Namibia	NA	250	9	1	260	2.5	10
50	Argentina	AR	3864	50	418	4332	44.9	10

Table 8

Top US MSAs with population of at least 250k, by developers per capita.

MSA Name	Count Contributors	Population	Contributors/100k
Boulder, CO	995	326196	305
San Jose-Sunnyvale-Santa Clara, CA	4587	1990660	230
San Francisco-Oakland-Hayward, CA	10702	4731803	226
Seattle-Tacoma-Bellevue, WA	8830	3979845	221
Ann Arbor, MI	600	367601	163
Champaign-Urbana, IL	365	226033	161
Austin-Round Rock, TX	3370	2227083	151
Durham-Chapel Hill, NC	759	644367	117
Portland-Vancouver-Hillsboro, OR-WA	2751	2492412	110
Charlottesville, VA	238	218615	108
Boston-Cambridge-Newton, MA-NH	5221	4873019	107
Santa Cruz-Watsonville, CA	249	273213	91
Denver-Aurora-Lakewood, CO	2555	2967239	86
Madison, WI	574	664865	86
Raleigh, NC	1159	1390785	83
Salt Lake City, UT	989	1232696	80

(continued on next page)

Table 8 (continued)

MSA Name	Count Contributors	Population	Contributors/100k
Lafayette-West Lafayette, IN	166	233002	71
Trenton, NJ	251	367430	68
Gainesville, FL	227	329128	68
Santa Maria-Santa Barbara, CA	291	446499	65
New York-Newark-Jersey City, NY-NJ-PA	11579	19216182	60
Provo-Orem, UT	390	648252	60
San Diego-Carlsbad, CA	1903	3338330	57
College Station-Bryan, TX	148	264728	55
Pittsburgh, PA	1185	2317600	51
Fort Collins, CO	181	356899	50
Nashville-Davidson-Murfreesboro-Franklin, TN	952	1934317	49
Burlington-South Burlington, VT	104	220411	47
Athens-Clarke County, GA	101	213750	47
Atlanta-Sandy Springs-Roswell, GA	2668	6020364	44
San Luis Obispo-Paso Robles-Arroyo Grande, CA	124	283111	43
Washington-Arlington-Alexandria, DC-VA-MD-WV	2698	6280487	42
Eugene, OR	161	382067	42
Minneapolis-St. Paul-Bloomington, MN-WI	1554	3640043	42
Bellingham, WA	96	229247	41
Los Angeles-Long Beach-Anaheim, CA	5533	13214799	41
Chicago-Naperville-Elgin, IL-IN-WI	3876	9458539	40
Lincoln, NE	132	336374	39
Boise City, ID	272	749202	36
Rochester, NY	376	1069644	35
Tucson, AZ	347	1047279	33
Santa Rosa, CA	154	494336	31
Orlando-Kissimmee-Sanford, FL	804	2608147	30
Philadelphia-Camden-Wilmington, PA-NJ-DE-MD	1883	6102434	30
Vallejo-Fairfield, CA	138	447643	30
Charlotte-Concord-Gastonia, NC-SC	771	2636883	29
Kansas City, MO-KS	637	2157990	29
Rochester, MN	66	221921	29
Columbus, OH	603	2122271	28
Fargo, ND-MN	67	246145	27

Table 9

The concentration of workers with higher education, workers in high-tech industries, and OSS contributors across NUTS2 regions of countries, measured by the Adjusted Geographic Concentration (AGC). The concentration of OSS contributors in particular regions is consistently higher than the concentration of both alternative populations.

Country	Country Code	$AGC_{TertEdu}$	AGC_{HiTech}	AGC_{OSS}
Portugal	PT	0.10	0.15	0.22
Italy	IT	0.07	0.19	0.22
France	FR	0.08	0.19	0.28
Belgium	BE	0.06	0.12	0.29
Greece	EL	0.09	0.31	0.29
Sweden	SE	0.05	0.20	0.30
Spain	ES	0.07	0.24	0.31
Ireland	IE	0.05	0.09	0.31
Croatia	HR	0.02	0.25	0.32
Netherlands	NL	0.06	0.10	0.32
Switzerland	CH	0.05	0.10	0.33
Finland	FI	0.05	0.17	0.36
Denmark	DK	0.11	0.28	0.38
Austria	AT	0.06	0.17	0.39
United Kingdom	UK	0.08	0.17	0.40
Poland	PL	0.08	0.22	0.41
Germany	DE	0.06	0.16	0.42
Norway	NO	0.06	0.24	0.44
Romania	RO	0.17	0.38	0.47
Bulgaria	BG	0.17	0.39	0.54
Slovakia	SK	0.10	0.15	0.59
Serbia	RS	0.19	0.40	0.61
Slovenia	SI	0.14	0.23	0.61
Lithuania	LT	0.14	0.52	0.64
Hungary	HU	0.18	0.26	0.66
Czechia	CZ	0.14	0.18	0.67

Table 10

Summary statistics table of key national statistics.

	OSS /mm	PPP GNI PC	HDI 2019	WVS Trust	IPI 2019	ECI	AI/DL papers	Internet Pen.
N Obs.	183	176	178	75	117	157	183	180
Mean	3.27	21.52	0.72	25.05	6.62	-0.00	-4.63	54.95
Stdev.	1.99	21.18	0.15	18.35	1.54	1.00	4.92	29.04
Min.	-1.20	0.79	0.39	2.10	2.35	-2.32	-9.21	2.00
25%	1.95	5.18	0.60	12.00	5.79	-0.73	-9.21	27.88
50%	3.34	13.64	0.74	20.60	6.52	-0.05	-9.21	58.50
75%	4.85	32.63	0.84	33.30	7.85	0.70	0.11	80.12
Max	7.54	92.27	0.96	73.90	9.61	2.27	4.22	98.30

Table 11

Summary statistics table of key NUTS2 statistics.

	OSS	GDP PC	Pop	PopDens	EVSTrust	EQI	R&D Spd	%HighTech	EEPatents	Patents	%TertEdu
N Obs.	297	285	297	288	218	149	259	279	259	259	293
Mean	1.00	4.40	6.13	4.97	0.38	-0.18	3.10	4.11	-0.14	1.47	33.93
Stdev.	0.49	0.25	0.35	1.20	0.20	0.99	1.32	2.25	1.41	1.57	10.17
Min	-0.79	3.61	4.48	1.19	0.00	-2.26	-1.45	0.70	-2.30	-2.70	11.80
25%	0.71	4.26	5.94	4.29	0.22	-0.95	2.29	2.50	-1.26	0.28	26.40
50%	1.03	4.46	6.16	4.86	0.33	-0.29	3.14	3.70	-0.18	1.76	33.10
75%	1.32	4.58	6.34	5.66	0.53	0.50	4.09	5.15	0.90	2.64	41.10
Max	2.24	4.98	7.09	8.91	0.81	2.32	6.48	12.90	3.37	4.53	59.86

Table 12

Spearman Correlation table of key national statistics.

	OSS /mm	PPP GNI PC	HDI 2019	WVS Trust	IPI 2019	ECI	AI/DL papers	Internet Pen.
OSS /mm	1.000	0.785	0.860	0.677	0.874	0.820	0.665	0.775
PPP GNI PC	0.785	1.000	0.963	0.698	0.846	0.823	0.710	0.909
HDI 2019	0.860	0.963	1.000	0.681	0.866	0.847	0.743	0.906
WVS Trust	0.677	0.698	0.681	1.000	0.659	0.619	0.602	0.642
IPI 2019	0.874	0.846	0.866	0.659	1.000	0.798	0.734	0.763
ECI	0.820	0.823	0.847	0.619	0.798	1.000	0.757	0.831
AI/DL papers	0.665	0.710	0.743	0.602	0.734	0.757	1.000	0.712
Internet Pen.	0.775	0.909	0.906	0.642	0.763	0.831	0.712	1.000

Table 13

Spearman Correlation table of NUTS2 statistics.

	OSS	GDP PC	Pop	PopDens	EVSTrust	EQI	R&D Spd	%HighTech	EEPatents	Patents	%TertEdu
OSS	1.000	0.567	0.249	0.355	0.471	0.455	0.368	0.605	0.542	0.542	0.664
GDP PC	0.567	1.000	0.025	0.321	0.675	0.690	0.791	0.441	0.743	0.837	0.579
Pop	0.249	0.025	1.000	0.406	-0.282	-0.202	-0.414	0.235	0.280	0.192	0.045
PopDens	0.355	0.321	0.406	1.000	-0.100	0.115	0.130	0.457	0.269	0.300	0.258
EVSTrust	0.471	0.675	-0.282	-0.100	1.000	0.658	0.655	0.253	0.528	0.588	0.546
EQI	0.455	0.690	-0.202	0.115	0.658	1.000	0.672	0.177	0.579	0.711	0.662
R&D Spd	0.368	0.791	-0.414	0.130	0.655	0.672	1.000	0.407	0.568	0.671	0.495
%HighTech	0.605	0.441	0.235	0.457	0.253	0.177	0.407	1.000	0.505	0.450	0.586
EEPatents	0.542	0.743	0.280	0.269	0.528	0.579	0.568	0.505	1.000	0.886	0.442
Patents	0.542	0.837	0.192	0.300	0.588	0.711	0.671	0.450	0.886	1.000	0.455
%TertEdu	0.664	0.579	0.045	0.258	0.546	0.662	0.495	0.586	0.442	0.455	1.000

Table 14

Variance Inflation Factors (VIF) for country-level regression models, models 1-5. Despite significant pairwise correlations (see correlation table), VIF scores remain at reasonable levels.

	M1	M2	M3	M4	M5
PPP GNI per Cap.	2.7	3.2	3.8	3.2	3.5
Internet Penetration	2.7	5.8	3.3	3.7	6.6
Population	1	1.1	1.1	1.1	1.2
HDI	N/A	6.8	N/A	N/A	8.3
IPI	N/A	N/A	3.4	N/A	N/A
ECI	N/A	N/A	N/A	3.4	3.7

Table 15

Variance Inflation Factors (VIF) for NUTS2-level regression models. Despite significant pairwise correlations (see correlation table), VIF scores remain at reasonable levels.

	M1	M2	M3	M4	M5	M6	M7
Internet Pen.	2.2	2.8	2.3	2.3	2.4	2.5	2.3
GDP PC	2.3	2.8	6.2	2.6	3.8	4.5	2.5
Pop	1.2	1.3	2.6	1.2	1.5	1.4	1.2
PopDens	1.4	1.5	1.4	1.5	1.5	1.5	1.4
EVSTrust	N/A	2.4	N/A	N/A	N/A	N/A	N/A
R&D Spd	N/A	N/A	6.3	N/A	N/A	N/A	N/A
%HighTech	N/A	N/A	N/A	1.5	N/A	N/A	N/A
EEPatents	N/A	N/A	N/A	N/A	2.8	N/A	N/A
Patents	N/A	N/A	N/A	N/A	N/A	4.2	N/A
%TertEdu	N/A	N/A	N/A	N/A	N/A	N/A	1.5

Table 16

OLS regression models relating EU NUTS2 counts of GitHub contributors per 100k inhabitants (log-transformed) and socio-economic indicators. We report (generic) heteroskedasticity robust standard errors. The models in this table do not consider the potential bias introduced by spatial autocorrelation. The model diagnostics in the following table suggest that these results may not be robust.

	Active GitHub Contributors/100k Inhab. European NUTS2 (log, 2021)						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Internet Penetration	0.010**	0.007	0.012**	0.011***	0.011**	0.011**	0.004
—(% of Pop. 2017)	(0.005)	(0.006)	(0.005)	(0.004)	(0.004)	(0.004)	(0.004)
GDP per Cap.	0.754***	0.694***	0.014	0.450***	0.274	0.342	0.426***
—(log Eur, 2017)	(0.147)	(0.167)	(0.262)	(0.152)	(0.181)	(0.211)	(0.155)
Population	0.270***	0.381***	0.486***	0.266***	0.169**	0.216**	0.285***
—(log, 2020)	(0.083)	(0.095)	(0.136)	(0.087)	(0.085)	(0.085)	(0.079)
Population Dens.	0.043*	0.052**	0.048	-0.019	0.064**	0.058*	0.027
—(log, 2017)	(0.026)	(0.026)	(0.029)	(0.023)	(0.029)	(0.030)	(0.023)
EVS Trust		0.415**					
—(2017)		(0.176)					
R&D Spend. per Cap.			0.128**				
—(log, 2017)			(0.052)				
% Empl. High-Tech				0.089***			
—(2019/20)				(0.011)			
Patents Elec-Eng./100k					0.072***		
—(log, 2017)					(0.023)		
Patents/100k						0.050*	
—(log, 2017)						(0.027)	
% with Tertiary Edu.							0.022***
—(2019/20)							(0.002)
Observations	276	198	258	262	258	258	276
Adjusted R ²	0.388	0.428	0.410	0.503	0.406	0.395	0.530
Residual Std. Error	0.371	0.354	0.347	0.328	0.348	0.352	0.325
F Statistic	36.8***	27.1***	34.5***	50.0***	33.0***	31.0***	52.8***

*p<0.1; **p<0.05; ***p<0.01

Spatial Regression Diagnostics

In this section we report linear regression models fit via OLS on the NUTS2 data in Table 17, as well as their spatial regression diagnostics. As suggested in the main text, we observe significant spatial autocorrelation of the key dependent variable (OSS contributors per 100k inhabitants) via Moran's I. We also examined spatial autocorrelation

among the residuals of linear regression models fit via OLS using Moran's I and a Lagrange Multiplier test (Anselin (1988); Cliff and Ord (1981)). We find significant spatial autocorrelation, motivating the use of the GMM model (Arraiz et al. (2010)) in the manuscript. We note that despite these (necessary) adjustments, the results (in terms of estimated effect sizes and statistical significance) are highly similar.

Table 17

OLS regression spatial diagnostics for the models in Table 17 relating EU NUTS2 counts of GitHub contributors per 100k inhabitants (log-transformed) and socio-economic indicators. These tests indicate significant spatial correlations, necessitating the use of the more sophisticated GMM model in the text to adjust for potential biases.

	Model 1	Model 2	Model 3	Model 4	Model 5	Model 6	Model 7
Moran's I (residual)	3.16	3.36	2.03	2.67	1.85	2.36	1.89
Moran's I p-value	0.00	0.00	0.04	0.01	0.06	0.02	0.06
Lagrange Multiple error	8.91	9.82	3.46	6.10	2.87	4.80	3.00
Lagrange Multiple error prob.	0.00	0.00	0.06	0.01	0.09	0.03	0.08

References

- Aksulu, A., Wade, M.R., 2010. A comprehensive review and synthesis of open source research. *Journal of the Association for Information Systems* 11 (11), 6.
- Albusays, K., Bjorn, P., Dabbish, L., Ford, D., Murphy-Hill, E., Serebrenik, A., Storey, M.-A., 2021. The diversity crisis in software development. *IEEE Software* 38 (2), 19–25.
- Alhusen, H., Bennat, T., Bizer, K., Cantner, U., Horstmann, E., Kalthaus, M., Proeger, T., Sternberg, R., Töpfer, S., 2021. A new measurement conception for the doing-using-interacting mode of innovation. *Res. Pol.* 50 (4).
- Andreessen, M., 2011. Why software is eating the world. *Wall Street Journal* 20 (2011), C2.
- Anselin, L., 1988. Lagrange multiplier test diagnostics for spatial dependence and spatial heterogeneity. *Geographical analysis* 20 (1), 1–17.
- Anthes, G., 2016. Open source software no longer optional. *Communications of the ACM* 59 (8), 15–17.
- Arraiz, I., Drukker, D.M., Kelejian, H.H., Prucha, I.R., 2010. A spatial cliff-ord-type model with heteroskedastic innovations: Small and large sample results. *Journal of Regional Science* 50 (2), 592–614.
- Balland, P.-A., Boschma, R., 2021. Mapping the potentials of regions in europe to contribute to new knowledge production in industry 4.0 technologies. *Regional Studies* 1–15.
- Bettencourt, L.A., Ostrom, A.L., Brown, S.W., Roundtree, R.I., 2002. Client co-production in knowledge-intensive business services. *California management review* 44 (4), 100–128.
- Blind, K., Böhm, M., Grzegorzewska, P., Katz, A., Muto, S., Pätsch, S., Schubert, T., 2021. The impact of open source software and hardware on technological independence, competitiveness and innovation in the eu economy.
- Boschma, R., 2005. Proximity and innovation: a critical assessment. *Regional studies* 39 (1), 61–74.
- Boschma, R.A., Ter Wal, A.L., 2007. Knowledge networks and innovative performance in an industrial district: the case of a footwear district in the south of italy. *Industry and innovation* 14 (2), 177–199.
- Branstetter, L.G., Drev, M., Kwon, N., 2019. Get with the program: Software-driven innovation in traditional manufacturing. *Management Science* 65 (2), 541–558.
- Chattergoon, B., Kerr, W.R., 2022. Winner takes all? tech clusters, population centers, and the spatial transformation of us invention. *Research Policy* 51 (2), 104418.
- Cliff, A.D., Ord, J.K., 1981. *Spatial processes: models & applications*. Taylor & Francis.
- Dahlander, L., Gann, D.M., Wallin, M.W., 2021. How open is innovation? a retrospective and ideas forward. *Research Policy* 50 (4), 104218.
- Dahlander, L., Wallin, M.W., 2006. A man on the inside: Unlocking communities as complementary assets. *Research policy* 35 (8), 1243–1259.
- Dahlberg, S., Sundström, A., Holmberg, S., Rothstein, B., Alvarado Pachon, N., Dalli, C. M., 2021. The Quality of Government Basic Dataset. University of Gothenburg: The Quality of Government Institute.
- Doehne, M., Rost, K., 2021. Long waves in the geography of innovation: The rise and decline of regional clusters of creativity over time. *Research Policy* 50 (9), 104298.
- Eghbal, N., 2020. *Working in public: the making and maintenance of open source software*. Stripe Press.
- Ellison, G., Glaeser, E.L., 1997. Geographic concentration in us manufacturing industries: a dartboard approach. *Journal of political economy* 105 (5), 889–927.
- Essletzbichler, J., 2015. Relatedness, industrial branching and technological cohesion in us metropolitan areas. *Regional Studies* 49 (5), 752–766.
- Etzkowit, H., Zhou, C., 2017. *The triple helix: University-industry-government innovation and entrepreneurship*. Routledge.
- Fackler, T., Laurentsyeve, N., 2020. Gravity in online collaborations: Evidence from github. *CESifo Forum*, Vol. 21. München: ifo Institut-Leibniz-Institut für Wirtschaftsforschung an der, pp. 15–20.
- Fershtman, C., Gandal, N., 2011. Direct and indirect knowledge spillovers: the social network of open-source projects. *The RAND Journal of Economics* 42 (1), 70–91.
- Florida, R., 2002. *The rise of the creative class*, Vol. 9. Basic books New York.
- Florida, R., Rodríguez-Pose, A., Storper, M., 2021. *Cities in a post-covid world*. Urban Studies.
- Forman, C., van Zeebroeck, N., 2019. Digital technology adoption and knowledge flows within firms: Can the internet overcome geographic and technological distance? *Research Policy* 48 (8), 103697.
- Frenken, K., Cefis, E., Stam, E., 2015. Industrial dynamics and clusters: a survey. *Regional studies* 49 (1), 10–27.
- Fry, T., Dey, T., Karnauch, A., Mockus, A., 2020. A dataset and an approach for identity resolution of 38 million author ids extracted from 2b git commits. *Proc. of the 17th Int. Conf. on Mining Software Repositories*.
- Furtado, L.B., Cartaxo, B., Treude, C., Pinto, G., 2020. How successful are open source contributions from countries with different levels of human development? *IEEE Software* 38 (2), 58–63.
- Gerosa, M., Wiese, I., Trinkenreich, B., Link, G., Robles, G., Treude, C., Steinmacher, I., Sarma, A., 2021. The shifting sands of motivation: Revisiting what drives contributors in open source. 2021 IEEE/ACM 43rd International Conference on Software Engineering (ICSE).
- GESIS Data Archive, 2017. *European values study 2017: Integrated dataset*. 10.4232/1.13560.
- Gomez-Lievano, A., Patterson-Lomba, O., 2021. Estimating the drivers of urban economic complexity and their connection to economic performance. *Royal Society Open Science* 8 (9), 210670.
- Gonzalez-Barahona, J.M., Robles, G., Andradás-Izquierdo, R., Ghosh, R.A., 2008. Geographic origin of libre software developers. *Information Economics and Policy* 20 (4), 356–363.
- Gousios, G., Spinellis, D., 2012. Ghtorrent: Github's data from a firehose. 2012 9th IEEE Working Conference on Mining Software Repositories (MSR). IEEE, pp. 12–21.
- Greenstein, S., Nagle, F., 2014. Digital dark matter and the economic contribution of Apache. *Res. Pol.* 43 (4).
- Grier, D.A., 2015. The tyranny of geography. *IEEE Annals of the History of Computing* 48 (02), 100.
- Haerpfer, C., Inglehart, R., Moreno, A., Welzel, C., Kizilova, K., Diez-Medrano, J., 2017. Round seven-country-pooled datafile. *World Values Survey*.
- HDI, 2019. Human Development Index. United Nation Development Program.
- Hecht, B., Hong, L., Suh, B., Chi, E.H., 2011. Tweets from justin bieber's heart: the dynamics of the location field in user profiles. *Proc. of the SIGCHI conference on Human Factors in Computing Systems*, pp. 237–246.
- Hegyi, F., Zhu, M., Janosov, M., 2021. Measuring the impact of urban innovation districts. Publications Office of the European Union. <https://doi.org/10.2760/11053>.
- Heller, B., Marschner, E., Rosenfeld, E., Heer, J., 2011. Visualizing collaboration and influence in the open-source software community. *Proceedings of the 8th working conference on mining software repositories*, pp. 223–226.
- Hidalgo, C.A., 2021. Economic complexity theory and applications. *Nature Reviews Physics* 1–22.
- Hu, T., Li, J., Zhou, H., Li, C., Holmes, E.C., Shi, W., 2021. Bioinformatics resources for sars-cov-2 discovery and surveillance. *Briefings in bioinformatics*.
- Jensen, M.B., Johnson, B., Lorenz, E., Lundvall, B.-Å., Lundvall, B., 2007. Forms of knowledge and modes of innovation. *The learning economy and the economics of hope* 155.
- Juhász, S., 2021. Spinoffs and tie formation in cluster knowledge networks. *Small Business Economics* 56 (4), 1385–1404.
- Juhász, S., Lengyel, B., 2018. Creation and persistence of ties in cluster knowledge networks. *Journal of Economic Geography* 18 (6), 1203–1226.
- Kalliamvakou, E., Gousios, G., Blincoe, K., Singer, L., German, D.M., Damian, D., 2014. The promises and perils of mining github. *Proceedings of the 11th working conference on mining software repositories*, pp. 92–101.
- Kaminski, J., Hopp, C., Tykrová, T., 2019. New technology assessment in entrepreneurial financing—does crowdfunding predict venture capital investments? *Technological Forecasting and Social Change* 139, 287–302.
- Klepper, S., 2010. The origin and growth of industry clusters: The making of silicon valley and detroit. *Journal of Urban Economics* 67 (1), 15–32.
- Klinger, J., Mateos-Garcia, J., Stathouloupoulos, K., 2021. Deep learning, deep change? mapping the evolution and geography of a general purpose technology. *Scientometrics* 1–33.
- Krugman, P., 1991. Increasing returns and economic geography. *Journal of political economy* 99 (3), 483–499.
- Kunzner, V., Lindner, T., Puck, J., 2019. Benefitting from immigration: The value of immigrants country knowledge for firm internationalization. *Journal of International Business Policy* 2 (4), 356–375.
- Lakhani, K.R., Von Hippel, E., 2004. *How open source software works: free user-to-user assistance*. Produktentwicklung mit virtuellen Communities. Springer, pp. 303–339.
- Latané, B., Liu, J.H., Nowak, A., Bonevanto, M., Zheng, L., 1995. Distance matters: Physical space and social impact. *Personality and Social Psychology Bulletin* 21 (8), 795–805.
- Lee, J.-A., 2006. Government policy toward open source software: The puzzles of neutrality and competition. *Knowledge, Technology & Policy* 18 (4), 113–141.
- Lerner, J., Tirole, J., 2002. Some simple economics of open source. *The journal of industrial economics* 50 (2), 197–234.
- Lima, A., Rossi, L., Musolesi, M., 2014. Coding together at scale: Github as a collaborative social network. Eighth international AAAI conference on weblogs and social media.
- Maraut, S., Derris, H., Webb, C., Spiezia, V., Guellec, D., 2008. *The oecd regpat database: a presentation*.
- May, A., Wachs, J., Hannák, A., 2019. Gender differences in participation and reward on stack overflow. *Empirical Software Engineering* 24 (4), 1997–2019.
- Mewes, L., Broekel, T., 2020. Technological complexity and economic growth of regions. *Research Policy* 104156.
- Montandon, J.E., Silva, L.L., Valente, M.T., 2019. Identifying experts in software libraries and frameworks among GitHub users. 2019 IEEE/ACM 16th Int. Conf. on Mining Software Repositories (MSR).
- Mungiu-Pippidi, A., Dadaşov, R., 2016. Measuring control of corruption by a new index of public integrity. *European Journal on Criminal Policy and Research* 22 (3), 415–438.
- Nagle, F., 2018. Learning by contributing: Gaining competitive advantage through contribution to crowdsourced public goods. *Organization Science* 29 (4), 569–587.
- Nagle, F., 2019. Government technology policy, social value, and national competitiveness. *Harvard Business School Strategy Unit Working Paper* (19-103).
- Nagle, F., 2019. Open source software and firm productivity. *Management Science* 65 (3), 1191–1215.
- Nagle, F., Wheeler, D., Lifshitz-Assaf, H., Ham, H., Hoffman, J., 2020. Report on the 2020 foss contributor survey. The Linux Foundation Core Infrastructure Initiative.
- Neffke, F.M., 2019. The value of complementary co-workers. *Science advances* 5 (12), eaax3370.
- Overney, C., Meinicke, J., Kästner, C., Vasilescu, B., 2020. How to not get rich: An empirical study of donations in open source. *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering*, pp. 1209–1221.
- Pietri, A., Spinellis, D., Zaccarioli, S., 2019. The Software Heritage graph dataset: public software development under one roof. 2019 IEEE/ACM 16th Int. Conf. on Mining Software Repositories. IEEE, pp. 138–142.

- Pintar, N., Scherngell, T., 2021. The complex nature of regional knowledge production: Evidence on European regions. *Res. Pol.*
- Prana, G.A.A., Ford, D., Rastogi, A., Lo, D., Purandare, R., Nagappan, N., 2021. Including everyone, everywhere: Understanding opportunities and challenges of geographic gender-inclusion in OSS. *IEEE Trans. Soft. Eng.*
- Rastogi, A., Nagappan, N., Gousios, G., van der Hoek, A., 2018. Relationship between geographical location and evaluation of developer contributions in GitHub. *Proceedings of the 12th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement*, pp. 1–8.
- Raymond, E., 1999. The cathedral and the bazaar. *Knowledge, Technology & Policy* 12 (3), 23–49.
- Rey, S.J., Anselin, L., 2010. Pysal: A python library of spatial analytical methods. *Handbook of applied spatial analysis*. Springer, pp. 175–193.
- Riedl, C., Seidel, V.P., 2018. Learning from mixed signals in online innovation communities. *Organization Science* 29 (6), 1010–1032.
- Rothstein, B., Uslander, E.M., 2005. All for all: Equality, corruption, and social trust. *World Pol.* 58, 41.
- Rovolis, A., Tragaki, A., 2006. Ethnic characteristics and geographical distribution of immigrants in Greece. *European Urban and Regional Studies* 13 (2), 99–111.
- Sahay, S., 2019. Free and open source software as global public goods? what are the distortions and how do we address them? *The Electronic Journal of Information Systems in Developing Countries* 85 (4), e12080.
- Saxenian, A., 1990. Regional networks and the resurgence of silicon valley. *California Management Review* 33 (1), 89–112.
- Secundo, G., Perez, S.E., Martinaitis, Z., Leitner, K.H., 2017. An intellectual capital framework to measure universities' third mission activities. *Technological Forecasting and Social Change* 123, 229–239.
- Spiezia, V., 2003. Measuring regional economies. *Statistics Directorate of the OECD*.
- Steinmacher, I., Balali, S., Trinkenreich, B., Guizani, M., Izquierdo-Cortazar, D., Cuevas Zambrano, G.G., Gerosa, M.A., Sarma, A., 2021. Being a mentor in open source projects. *Journal of Internet Services and Applications* 12 (1), 1–33.
- Storper, M., Venables, A.J., 2004. Buzz: face-to-face contact and the urban economy. *Journal of economic geography* 4 (4), 351–370.
- Su, Y.-S., Hung, L.-C., 2009. Spontaneous vs. policy-driven: The origin and evolution of the biotechnology cluster. *Technological Forecasting and Social Change* 76 (5), 608–619.
- Takhteyev, Y., 2012. Coding places: Software practice in a South American city. MIT Press.
- Takhteyev, Y., Hilt, A., 2010. Investigating the geography of open source software through github. Manuscript.
- The International Telecommunication Union, 2019. Internet penetration.
- Tóth, G., Juhász, S., Elekes, Z., Lengyel, B., 2021. Repeated collaboration of inventors across European regions. *European Planning Studies* 1–21.
- Trippl, M., Tödtling, F., Lengauer, L., 2009. Knowledge sourcing beyond buzz and pipelines: evidence from the Vienna software sector. *Economic geography* 85 (4), 443–462.
- Trujillo, M. Z., Hbert-Dufresne, L., Bagrow, J. P., 2021. The penumbra of open source: projects outside of centralized platforms are longer maintained, more academic and more collaborative. 2106.15611.
- Uyarra, E., Ramlogan, R., 2016. The impact of cluster policy on innovation. *Handbook of innovation policy impact*. Edward Elgar Publishing.
- Valiev, M., Vasilescu, B., Herbsleb, J., 2018. Ecosystem-level determinants of sustained activity in open-source projects: A case study of the pypi ecosystem. *Proc. of the 2018 26th ACM Joint Meeting ESEC/FSE*.
- Wachs, J., Vedres, B., 2021. Does crowdfunding really foster innovation? evidence from the board game industry. *Technological Forecasting and Social Change* 168. <https://doi.org/10.1016/j.techfore.2021.120747>.
- Wagner, C., Strohmaier, M., Olteanu, A., Kiciman, E., Contractor, N., Eliassi-Rad, T., 2021. Measuring algorithmically infused societies. *Nature*.
- Wanzenböck, I., Neuländtner, M., Scherngell, T., 2020. Impacts of EU funded R&D networks on the generation of key enabling technologies: Empirical evidence from a regional perspective. *Papers in Regional Science* 99 (1), 3–24.
- Wanzenböck, I., Wesseling, J.H., Frenken, K., Hekkert, M.P., Weber, K.M., 2020. A framework for mission-oriented innovation policy: Alternative pathways through the problem–solution space. *Sci. and Pub. Pol.* 47 (4).
- Watney, C., 2020. Clusters rule everything around me. *Works in progress* (2).
- Weterings, A., Boschma, R., 2006. The impact of geography on the innovative productivity of software firms in the Netherlands. *Regional Development in the Knowledge Economy*. Routledge, London and New York.
- Wright, N., Nagle, F., Greenstein, S.M., 2020. Open source software and global entrepreneurship. *Harvard Business School Technology & Operations Mgt. Unit Working Paper* (20-139), 20–139.
- Wu, L., Jin, F., Hitt, L.M., 2018. Are all spillovers created equal? a network perspective on information technology labor movements. *Management Science* 64 (7), 3168–3186.
- Johannes Wachs** is an assistant professor at the Vienna University of Economics and Business and a faculty member at the Complexity Science Hub Vienna. He researches the social and economic impact of the web and digital technologies, using empirical methods. In particular, he studies the growing socio-economic importance of Web-based forms of collaboration such as crowdfunding and open source software. He has a PhD in Network Science from Central European University and has previously been affiliated with RWTH Aachen University and the Oxford Internet Institute.
- Mariusz Nitecki** is a research assistant at the Vienna University of Economics and Business and a master's student in Data Science at the Vienna University of Technology, specializing in machine learning and high performance computing. He has a bachelor's degree in Information Systems, and wrote his thesis on algorithmic trading systems using artificial neural networks. His primary field of interest is the application of machine learning in finance.
- William Schueller** is a postdoc at the Complexity Science Hub Vienna and the Medical University of Vienna. His current projects focus on computing systemic risk indicators in different contexts, such as food supply chains or OSS contribution networks. After a master in physics and complex systems at the ENS Lyon, and two years teaching maths and physics in Galatasaray University in Istanbul, he obtained a PhD in computer science from INRIA Bordeaux Sud-Ouest on multi-agent models of language evolution.
- Axel Polleres** heads the Institute for Data, Process and Knowledge Management of Vienna University of Economics and Business (WU Wien), which he joined in September 2013 as a full professor in the area of "Data and Knowledge Engineering". He is also a faculty member of the Complexity Science Hub Vienna and has been visiting professor at Stanford University in 2018. He obtained his Ph.D. and habilitation from Vienna University of Technology. His research focuses on ontologies, query languages, logic programming, configuration technologies, Artificial Intelligence, Semantic Web, Linked Open Data, Knowledge Graphs and their applications for Knowledge Management. Moreover, he actively contributed to international standardisation efforts within the World Wide Web Consortium (W3C) where he co-chaired the W3C SPARQL working group.